What Korea Says

A Family Feud–style survey quiz. NVIDIA Nemotron on Nebius Token Factory writes the first answer board; real players turn it into Korean public opinion.

Play: nebius.woogamer.win/play/?lang=en · Code: github.com/woogamer/what-korea-says (MIT)

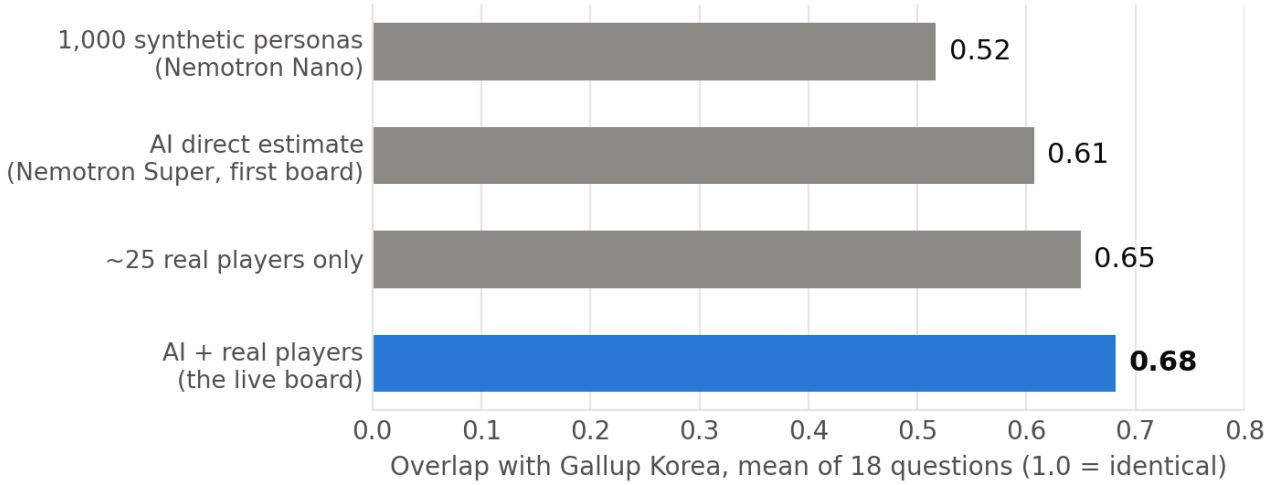
126
real players in two days
(link shared in group chats, no ads)

887
rounds played
9.5 questions per player

11
players finished
all 30 questions

0.61 → 0.68
overlap with Gallup Korea
AI alone → AI + players

Real people beat synthetic people, and AI + people beats both

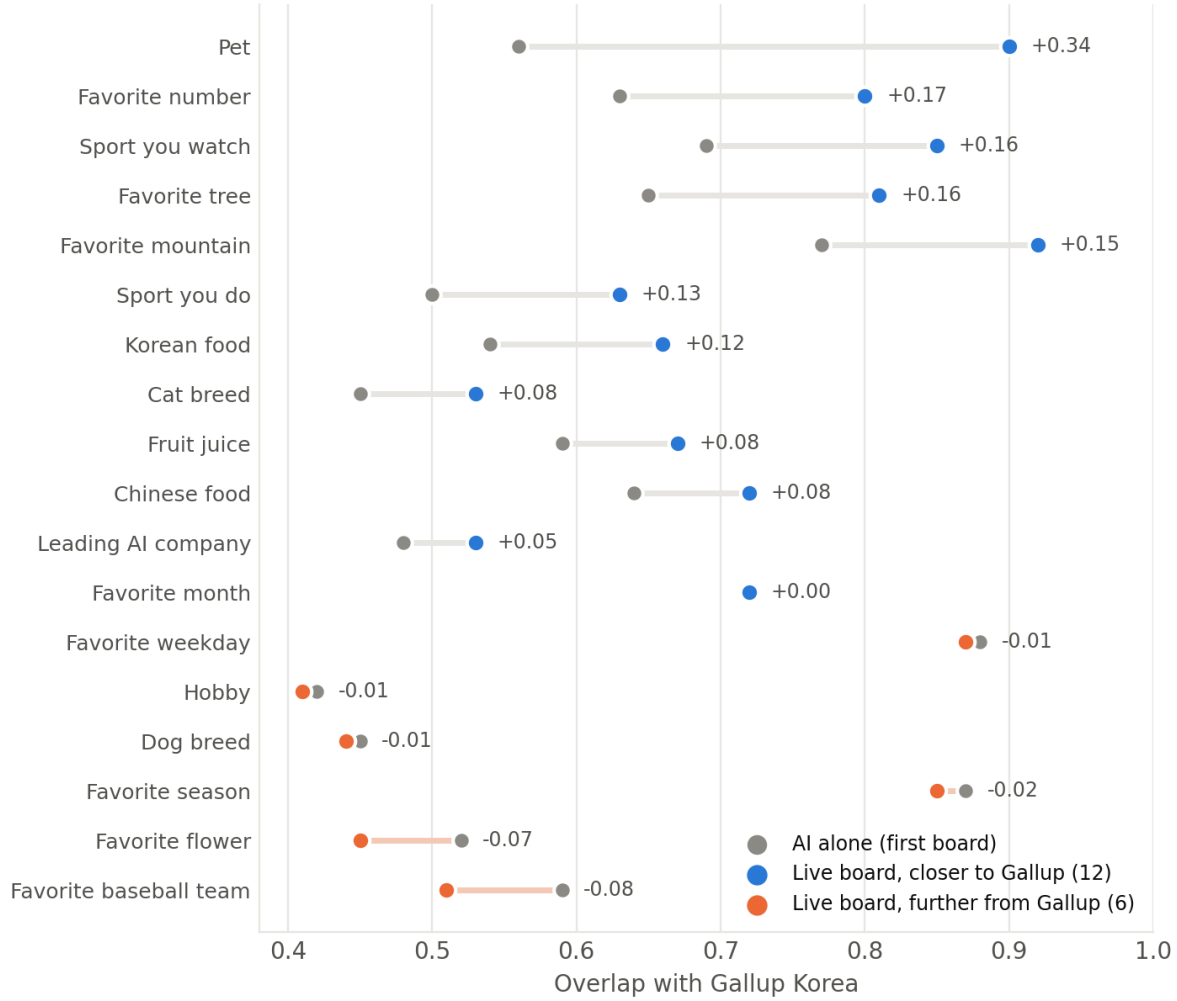


Overlap = Σ min(board %, Gallup %) / Σ Gallup %, averaged over the 18 questions that Gallup Korea also asked (2024 face-to-face survey of 1,777 people, and two 2026 releases).

- 1,000 synthetic personas (0.52) lose to about 25 real players per question (0.65). Synthetic people all make the same mistake.
- Nemotron's direct estimate (0.61) is a good first draft. Each real answer updates it with a Bayesian rule in which the AI counts as about 8 people.
- The live board (0.68) beats both sources alone, as the simulation predicted.

Live results by question

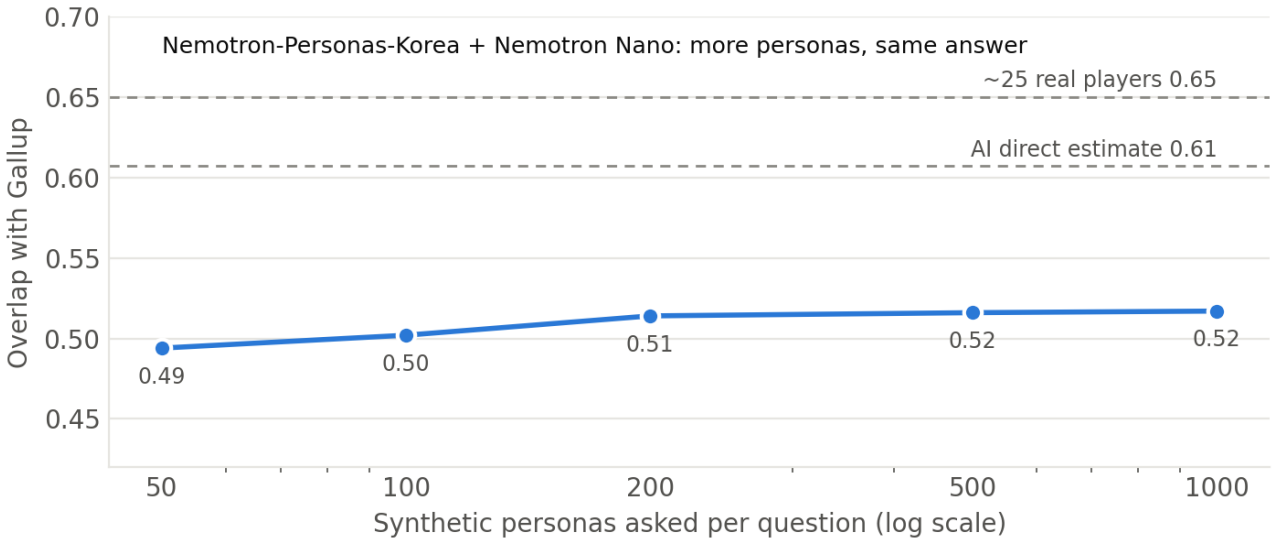
Each row is one Gallup question: the grey dot is Nemotron's first board, the coloured dot is the board after 22–32 Korea-resident players. The board moved closer to Gallup on 12 questions and further on 6.



- Biggest gains are where the AI's first guess was furthest from Koreans: pet, favourite number, sport watched, tree, mountain.
- Losses come from small, self-selected samples: a few fans of one baseball team move a 25-person board.

Why not just simulate people? More personas did not help

We asked 1,000 Nemotron-Personas-Korea profiles all 18 Gallup questions with Nemotron 3 Nano on Nebius (about 18,000 calls, $1.13), then scored random subsets.



The error is bias, not sampling noise. A thousand different people converge on one answer, and none of them says “none”:

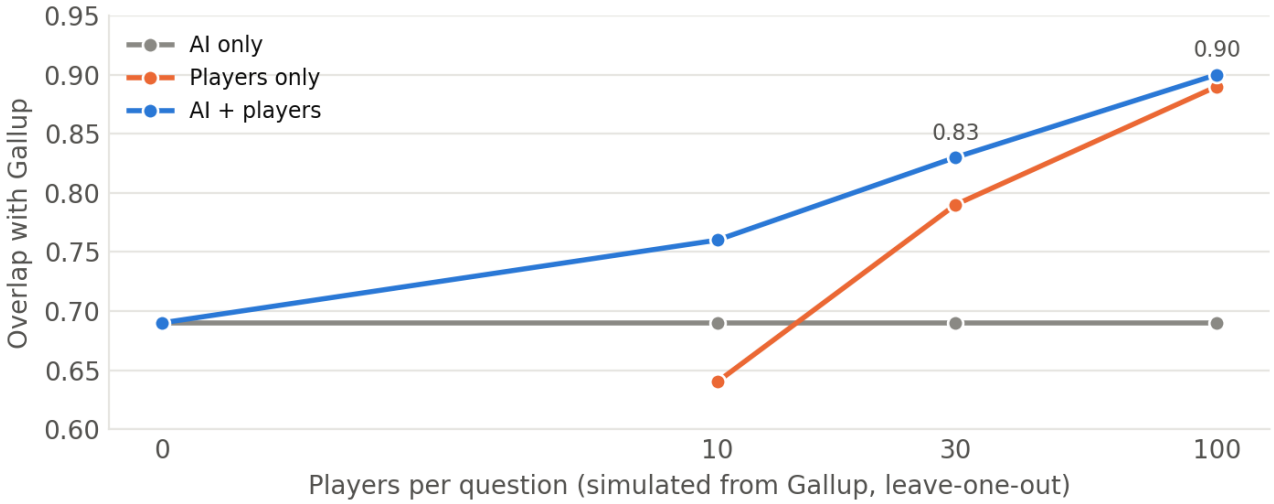
Question	Gallup Korea	1,000 personas
Favourite tree	pine 36%	cherry 84%
Leading AI company	Samsung 37%	Google 82%
Favourite number	7 at 34%	7 at 83%
Favourite weekday	Friday 34%, Saturday 33%	Saturday 75%
Favourite baseball team	none 46%	one team 70%

Asking Nemotron 3 Super directly for the national distribution scored higher on every held-out Gallup set (2024: 0.66 vs 0.51; 2026, released after training: 0.62 vs 0.26), so the game starts from the direct estimate.

How the board learns

$$\text{board}_k = (\kappa \cdot \text{AI}_k + \text{players}_k) / (\kappa + n)$$

- κ ≈ 8, fit by empirical Bayes on 30 Gallup question–survey pairs (2019, 2024, 2026): the AI is worth about 8 real respondents.
- New answers the AI missed join the board once two players give them. Only players living in Korea count, one vote per browser per question.
- In simulation (players drawn from Gallup, leave-one-out), AI + players beat either source alone at every size:



What did not work

- Age reweighting. Players skew to their 30s–40s (20s 13, 30s 36, 40s 30, 50s 12, 60+ 10). Reweighting to Korea's adult population lowered the live board from 0.68 to 0.63: a few older players got very large weights, and they did not answer more like Gallup than younger ones.
- A global bias correction learned across questions gained only +0.02. The AI is overconfident: its 80% intervals covered the truth 65% of the time.

Built with

Nebius Token Factory · NVIDIA Nemotron 3 Super (first boards, Gallup mapping, live answer judge: fast with reasoning off on screen, careful with reasoning on in the background) · Nemotron 3 Nano + Nemotron-Personas-Korea (the rejected baseline) · FastAPI · SQLite · JavaScript. Reproduce: **python -m game.live_report**, **python -m game.validate_persona_n**, **python -m game.evolve**.

Limits: players are self-selected friends and family; AI boards are model estimates, not survey results; Gallup publishes only its top answers, so it is a reference point, not an answer key.